


CMED → AIMS LAB — Software Integration

Document 1 of 2 — for CMED's software, database and data team · The APIs, the keys, the fields and the data model · AIMScribe v3 · AIMS LAB, United International University · 9 October 2026

	The two documents
1. This one	Software integration — the four APIs, the key, the exact endpoints, every field, every reply code, where the data lands, and how to test it. For the team writing the code
2. The other	Server deployment — the machine, the region, DNS, TLS, ports, the services, health checks and monitoring. For the team helping us deploy

What this integration is, in one paragraph

AIMScribe records the consultation. It runs as a small program on **the same laptop** as the doctor's browser. CMED's page tells that program when a patient is opened and when the prescription is built, so it knows where one consultation ends and the next begins. Separately, **CMED's server** sends us the patient's information and the prescription, so the recording has a clinical record attached. **Audio never goes to CMED, and nothing clinical is ever served back to CMED.**



Audio never reaches CMED. Nothing clinical is ever served back to CMED.

The two downward paths never meet: the recorder sends audio to us, and CMED's server sends us the clinical record. Neither party ever receives the other's data, and there is no endpoint by which it could. **CMED's page never learns the AIMS LAB address either - only CMED's own server holds it.**

Every rule here is taken from the program that validates these messages on arrival. **If this document and the software disagree, the software is right and this is a bug** — please tell us.

Contents

Part	Covers	Sections
A	The keys and the exact endpoints — start here	1–3
B	The five fields, and how the four messages relate	4–5
C	The two browser signals — API 1 and API 3 Channel A	6–11

D	The two server messages — API 2 and API 3 Channel B	12–17
E	Previous prescriptions — today, and what we will add	18–19
F	The data model — where every field lands	20–21
G	Testing, handover, and one page to remember	22–25

The keys and the exact endpoints

Copy these. Everything after this explains them

1 · The gateway

```
BASE_URL = https://<aims-host>          <-- AIMS LAB sends this with the key
```

```
Transport  HTTPS only, TLS 1.2 or 1.3, port 443. HTTP redirects to HTTPS
Encoding   JSON, UTF-8
Body limit 1 MB
Auth       X-CMED-Key: <key>          on the two clinical endpoints only
```

Hold the host in configuration, not compiled in. You will get a staging host first and a production host later, and they differ.

2 · The authentication, in full

	Channel B — your server to ours	Channel A — your page to the recorder
Credential	X-CMED-Key header	none — and none is possible
Why	Server to server, so a secret is safe	A secret in a web page is readable by anyone who opens the page
What authorises you	the key	your page's exact web address , on our allow-list, plus loopback-only, plus same-machine
Issued by	AIMS LAB, once, out of band. Shown once and not recoverable	you send us the address; we register it
Scope	one key for CMED, covering all seven clinics	one entry per address
Rotation	ask us — new key issued before the old is revoked, so there is no window where both fail	n/a
Where it must live	your server's secret store or environment	n/a

The key must never reach a browser — not in page source, not in a bundled script, not in a mobile app. A wrong or missing key returns **401 INVALID_KEY**, which is a hard stop: do not retry it, alert your team and contact us.

The address allow-list is a blocking prerequisite, and it is exact

Send AIMS LAB **every web address your page will be served from**, production and staging and preview. An address not on the list is refused with `ORIGIN_NOT_ALLOWED` and the connection never opens, so **nothing works at all**.

It must match character for character: `https://ehr.aaloclinic.com` and `https://www.ehr.aaloclinic.com` are different, so are `http` and `https`, and so is any non-default port.

3 · Every address CMED calls

#	Method	Full address	Auth	Body	OK
1	GET	https://<aims-host>/health	none	none	200
2	GET	https://<aims-host>/	none	none	200
3	POST	https://<aims-host>/api/v2/clinical/patient-information	key	JSON	202
4	POST	https://<aims-host>/api/v2/clinical/prescription	key	JSON	202

Protocol	Full address	Auth	Who
WebSocket	ws://127.0.0.1:5050/ws	address allow-list	your page, in the doctor's browser

Message on that socket	Which API	Sent when
{"command": "start", "trigger": {...}}	API 1	the doctor opens a patient
{"command": "prescription_built", ...}	API 3, Channel A	the doctor presses Build Prescription
{"command": "doctors"}	—	on page load — returns this PC's clinic code
{"command": "status"}	—	optional; the recorder also pushes status unprompted

All six, one line each, to copy

```
GET https://<aims-host>/health
GET https://<aims-host>/
POST https://<aims-host>/api/v2/clinical/patient-information
POST https://<aims-host>/api/v2/clinical/prescription
WS ws://127.0.0.1:5050/ws → "start"
WS ws://127.0.0.1:5050/ws → "prescription_built"
```

There is no GET for clinical data and no DELETE anywhere. CMED writes to us and never reads back — see section 18, which is the one place that will change. **Everything else under /api/v2/ is closed to CMED's key;** those endpoints belong to the recorders and use client certificates.

The five fields

The one rule that decides whether the whole integration works


4 · Three APIs, four messages

	What	From	When	Part
API 1	Start recording	the browser	the doctor opens a patient	C
API 2	Patient information	your server	the same moment	D
API 3 Channel A	Prescription built	the browser	the doctor presses Build Prescription	C
API 3 Channel B	The prescription	your server	the same moment	D

They go in pairs. Two moments, four messages. **API 3 is one event told to two places** — the recorder, so it knows the consultation has reached its end, and our server, so we receive the prescription. The earlier integration guide calls these 6a and 6b; this document says Channel A and Channel B. Three APIs, four messages.

4.1 · One whole consultation, start to finish

Read this once and the rest of the document is detail. Time runs down the page; each column is one party.



Three things to take from it. The four messages happen at only **two moments**, in pairs. **The match** — the step CMED cannot see — is what decides whether a recording is kept. And the recorder **keeps recording through the prescription**, closing only when the next patient is opened.

5 · The five fields

All four messages about one consultation carry these same five fields, with **exactly the same values, character for character**.


```
{
  "patient_id": "P0012345",
  "doctor_id": "DR0042",
  "hospital_id": "AALO_DH0LPUR",
  "start_time": "2026-10-09T10:14:32+06:00",
  "date": "2026-10-09"
}
```

Field	Type	Rule
patient_id	text	Letters, digits, hyphen, underscore. 1–64 characters. No spaces
doctor_id	text	same rule
hospital_id	text	same rule. Read it from the recorder — section 10
start_time	text	Date and time with the time-zone offset , e.g. +06:00
date	text	YYYY-MM-DD, the clinic's local date

Three mistakes that break the match

- 1. Sending the prescription's own time as start_time.** Both halves of API 3 repeat the **original** start_time from when the patient was opened. This is the most common integration mistake, and it is silent: both messages succeed and the prescription simply never finds its recording.
- 2. Leaving off the offset.** 2026-10-09T10:14:32 is rejected.
- 3. A space or other character in an identifier.** These become folder names on disk.

5.1 One object, four messages



Pass the same object to all four. Never rebuild it and never recompute the time.

In API 3a and API 3b, start_time is still the moment the PATIENT WAS OPENED - not the moment the prescription was built. That single mistake is the most common way this integration fails, and it is silent: every message succeeds, and the prescription simply never finds its recording.

Take hospital_id from the recorder's doctors command, not from your own settings.

Build the five fields once, into one object, when the doctor opens the patient. Keep it for the whole consultation and pass the same object to all four messages. Never rebuild them and never recompute the time. Do that and the match cannot fail; section 23 shows it in code.

The two browser signals

API 1 and API 3 Channel A - and the gate

6 · How the page connects

	Detail
Address	ws://127.0.0.1:5050/ws — the same on every clinic PC
Why ws:// and not wss://	It never leaves the machine. Loopback needs no certificate and the browser permits it
Protocol	One WebSocket held open for the clinic session. JSON, UTF-8
Largest message	64 KB
When to connect	when the page loads; reconnect quietly if it drops
Firewall changes	none, on CMED's side or the clinic's

7 · Every message and reply

CMED sends:

```
{ "command": "start", "request_id": "abc-123", ... }
```

Exactly one reply per message:

```
{
  "event":      "ack",                <-- "ack" below 400, "error" from 400
  "command":    "start",
  "status":     200,
  "code":       "RECORDING_STARTED", <-- ACT ON THIS
  "message":    "Recording started.", <-- for a human; may be reworded
  "request_id": "abc-123",
  "data":       { ... }
}
```

Act on code, never on message. The code is a promise; the wording is not.

8 · API 1 — a patient has been opened

```
{
  "command": "start",
  "request_id": "open-P0012345-1",
  "trigger": {
    "patient_id": "P0012345", "doctor_id": "DR0042",
    "hospital_id": "AALO_DHOLPUR",
    "start_time": "2026-10-09T10:14:32+06:00", "date": "2026-10-09"
  }
}
```

The microphone starts **immediately**, before any permission check finishes, so nothing of the consultation is missed. What comes back:

```
"data": {
  "session_id": "01M3MK4J4XEDG5NDKE4H5B5XA2", <-- KEEP THIS
  "started_at": "2026-10-09T10:14:32.865Z",
  "armed": false,
  "previous_session_stopped": true,
  "previous_session_id": "01M3MEYY...",
  "authorisation": "granted",
  "confirmation": "confirming"
}
```

9 · Every reply to API 1

Code	Meaning	What CMED should do
RECORDING_STARTED (200)	Recording, permission confirmed	Store the session_id. Show it is recording
RECORDING_PROVISIONAL (202)	Recording. Permission still being checked	Treat as success. Normal, not a warning
SESSION_ALREADY_ACTIVE (409)	This patient is already being recorded	Nothing; you probably sent it twice
GATE_NOT_ARMED (409)	The previous consultation has no prescription yet	Ask the doctor to build the previous prescription
AUTHORISATION_FAILED (401)	Refused by the server; recording stopped and deleted	Show it is not recording. Contact AIMS LAB
CLINIC_MISMATCH (401)	hospital_id is not this PC's clinic	Stop. Use section 10 to avoid this entirely
DOCTOR_NOT_AT_CLINIC (404)	That doctor is not registered here	Check doctor_id
MISSING_FIELD (400)	A field is missing or malformed	Fix and resend
INVALID_IDENTIFIER (400)	An id has disallowed characters	Fix and resend
DEVICE_NOT_ENROLLED (423)	This PC is not registered with AIMS LAB	Stop. Contact AIMS LAB
ORIGIN_NOT_ALLOWED (403)	Your page's address is not on the allow-list	Send us the address. Section 2
AGENT_NOT_READY (503)	The recorder is starting up	Wait and retry

RECORDING_PROVISIONAL is the one people get wrong. The microphone is live and audio is being written; the server has simply not finished confirming permission, which normally takes a fraction of a second. Show the same "recording" state as for RECORDING_STARTED. Do not warn and do not retry.

But a provisional start can still fail afterwards. If the server then refuses it you receive AUTHORISATION_FAILED or DEVICE_NOT_ENROLLED **a moment later, and the recording is stopped and deleted.** Do not treat the first reply as final: keep listening, and change the display if one of those arrives. The status messages in section 11 are the reliable way to show this.

10 · Never hardcode the clinic code — ask for it

Every AIMScribe laptop is enrolled to exactly one clinic, and a mismatch is **refused**. **So do not store clinic codes in your page — read them off the laptop on page load.**

```
send: { "command": "doctors", "request_id": "who-is-this-pc" }
```

```
reply: "data": {
  "hospital_id": "AALO_DHOLPUR",      <-- use THIS as hospital_id
  "doctors": [
    { "doctor_id": "DR0042", "full_name": "Dr Rahman" },
    { "doctor_id": "DR0051", "full_name": "Dr Akter" }
  ]
}
```

What comes back	Meaning	What the page should do
a code, and doctors listed	Normal	Proceed. Use the code; fill the doctor selector
a code, doctors: []	Also normal — a new clinic, or our server briefly unreachable	Proceed. Let the doctor be typed. An empty list must never block a clinic
hospital_id: null	This PC is not enrolled	Do not send API 1. Say so and contact us

This makes CLINIC_MISMATCH impossible rather than something to handle, and lets us rename a clinic without CMED changing anything. It already caused one failure in testing: a page was sending an old code after a clinic was renamed, and every recording was refused until the page was rebuilt. **Re-send doctors whenever the socket reopens.**

11 · API 3 Channel A, the gate, and status

```
{
  "command": "prescription_built",
  "request_id": "built-P0012345-1",
  "patient_id": "P0012345",
  "session_id": "01M3MK4J4XEDG5NDKE4H5B5XA2"    <-- from API 1
}
```

Field	Required	Rule
command	yes	prescription_built
patient_id	yes	the patient currently being recorded
session_id	yes	the one API 1 returned
occurred_at	no	date and time with offset


Code	Meaning	Do
GATE_ARMED (200)	Accepted	nothing
GATE_ALREADY_ARMED (200)	Sent twice	nothing; harmless
PATIENT_MISMATCH (409)	Not the patient being recorded	check the session
NO_ACTIVE_SESSION (409)	Nothing is recording	the consultation already ended

Building the prescription does NOT stop the recording

The doctor builds it, prints it, hands it over — and then **keeps talking**. That counselling is some of the most valuable audio in the study.

So prescription_built ends nothing. It **arms a gate**: this consultation has reached its end, so the next patient is allowed to close it. **The recording stops when the next patient is opened** — your next API 1, typically twenty to thirty seconds later.

11.1 What the recorder is doing, state by state



The third state is the one that gets missed.

prescription_built does not stop the recording. It moves the recorder from RECORDING to RECORDING + armed, which means "this consultation may now be closed by the next one". Audio keeps being written throughout.

If the next API 1 arrives without the gate being armed, it is refused with GATE_NOT_ARMED.

The doctor does	CMED sends	The recorder does
Opens patient A	API 1 + API 2	starts recording A
Consults	—	recording
Presses Build Prescription	API 3, both channels	keeps recording. Arms the gate

Prints, hands over, counsels	—	still recording — the valuable part
Opens patient B	API 1 + API 2	closes A, starts B

If the doctor finishes for the day or goes to lunch, they press Stop on the recorder's own on-screen control. CMED sends nothing, and must not try to. The socket also accepts stop, pause and resume, but **these are not CMED's to use**: each records a clinical reason that becomes part of the permanent record, including "the patient did not consent", which erases the recording everywhere. **Send only start and prescription_built.**

Status, pushed to you unprompted

```
{
  "event": "status",
  "state": "recording",           <-- idle | recording | paused | closing
  "is_recording": true, "is_paused": false,
  "session_id": "01M3MK4J...", "patient_ref": "P0012345",
  "armed": false,                <-- has the prescription signal arrived?
  "authorisation": "granted", "confirmation": "confirmed",
  "duration_seconds": 182.4,
  "level": 0.42,                 <-- 0.0 to 1.0, how loud the room is
  "connected_clients": 1
}
```

Show level as a small meter. A microphone that is unplugged, muted in Windows, or set to the wrong input records perfect silence and **nothing on screen looks wrong**. Without a meter that is discovered when somebody plays the file back, possibly weeks and hundreds of consultations later. It is the cheapest safeguard in the whole integration.

The two server messages

API 2 and API 3 Channel B - every field, every limit, every reply

12 · API 2 — patient information

POST https://<aims-host>/api/v2/clinical/patient-information

X-CMED-Key: <key>

```
{
  "patient_id": "P0012345", "doctor_id": "DR0042",
  "hospital_id": "AALO_DHOLPUR",
  "start_time": "2026-10-09T10:14:32+06:00", "date": "2026-10-09",

  "demographics": {
    "name": "Rahima Khatun", "sex": "female", "age_years": 34,
    "date_of_birth": "1992-03-11",
    "phone": "01700000000", "address": "Dholpur, Dhaka"
  },

  "paramedic": {
    "recorded_at": "2026-10-09T10:02:10+06:00",
    "weight_kg": 58.4, "height_cm": 157,
    "blood_pressure": "120/80", "pulse_bpm": 78,
    "temperature_c": 36.8, "spo2_percent": 98,
    "notes": "complains of headache"
  },

  "previous_visit": null,
  "female_details": {}, "male_details": {}
}
```

12.1 demographics — required

Field	Type	Accepted	If wrong	Lands in
name	text	up to 200. full_name also accepted	shortened	patients.full_name
sex	text	exactly female or male	whole message quarantined	patients.sex
age_years	number	0 to 130	stored empty	encounter_demographics
date_of_birth	text	YYYY-MM-DD	stored empty	patients.date_of_birth
phone	text	up to 40	shortened	patients.phone
address	text	up to 500	shortened	patients.address

demographics must be present on every visit, first or fiftieth. **sex is the one field that can fail the whole message** — send the word, not a code like M or F.

12.2 paramedic — optional, but send it when you have it

Field	Type	Accepted	Lands in
recorded_at	text	date and time with offset	paramedic_observations
weight_kg	number	0.5 to 499	weight_kg
height_cm	number	20 to 299	height_cm
blood_pressure	text	"120/80" — upper 30–300, lower 10–250	blood_pressure, also split into systolic and diastolic
pulse_bpm	whole number	10 to 300	pulse_bpm
temperature_c	number	25 to 45	temperature_c
spo2_percent	whole number	0 to 100	spo2_percent
notes	text	up to 2000	notes

A reading outside its range is stored empty — it never rejects the message. One mistyped temperature must not stop a prescription being saved.

Anything else inside paramedic is kept. If you measure something we have not listed — respiration rate, blood sugar, anything — send it. It is stored in an other field alongside the rest. **You do not need our permission and we do not need to change anything first.**

13 · API 3 Channel B — the prescription

POST <https://<aims-host>/api/v2/clinical/prescription>

X-CMED-Key: <key>

```
{
  "patient_id": "P0012345", "doctor_id": "DR0042",
  "hospital_id": "AALO_DHOLPUR",
  "start_time": "2026-10-09T10:14:32+06:00",    <-- the ORIGINAL
  "date": "2026-10-09",

  "issued_at": "2026-10-09T10:26:55+06:00",
  "items": [
    { "drug": "Napa 500 mg", "dose": "1 tablet",
      "frequency": "1+1+1", "duration": "5 days",
      "instructions": "after food" },
    { "drug": "Omeprazole 20 mg", "dose": "1 capsule",
      "frequency": "1+0+0", "duration": "14 days" }
  ],
  "diagnoses": ["Acute gastritis"],
  "investigations": ["CBC", "HbA1c"],
  "advice": "avoid spicy food",
  "follow_up": "2026-11-09",
  "notes": "patient counselled about diet"
}
```

Field	Type	Required	Rule	Lands in
issued_at	text	YES	date and time with offset	prescriptions.issued_at
items	list	YES	one per medicine. Send [] if none	prescription_items
items[].drug	text	YES	up to 300. Every item needs one	drug
items[].dose	text	no	up to 200	dose
items[].frequency	text	no	up to 200, e.g. 1+1+1	frequency
items[].duration	text	no	up to 200, e.g. 5 days	duration
items[].instructions	text	no	up to 1000	instructions
diagnoses	list of text	YES	Send [] if none	diagnoses

investigations	list of text	YES	Send [] if none	investigations
advice	text	no	up to 2000	prescriptions.advice
follow_up	text	no	YYYY-MM-DD if sent	follow_up_date
notes	text	no	up to 2000	prescriptions.notes

An empty list and a missing field are not the same thing. "diagnoses": [] means "the doctor recorded none"; leaving it out means "CMED did not send it", and that is reported as a problem. **Always send the list.** The order of items is preserved as the line number on the prescription.

14 · Corrections, repeats and retries

Situation	What to do	What happens
The doctor edits the prescription after sending	Send API 3 Channel B again with the same five fields	Stored as a new version. Both kept; the newest is current
Your network failed and you are unsure	Send it again, unchanged	Completely safe. 202 ALREADY_RECEIVED and one row, not two
You sent something wrong	Send the corrected version, same five fields	Same as an edit
You sent a field we do not recognise	nothing	Kept with the rest. Unknown fields are never an error

Every message is stored word for word, exactly as sent, before anything is read out of it. Nothing is ever lost because we failed to understand a field — if we later realise we should have been reading something, we can go back through the stored messages and get it.

15 · Every reply from the server

Code	Means	What CMED should do
202 ACCEPTED	Stored safely	Nothing. Mark it sent
202 ALREADY_RECEIVED	This exact message was already stored	Nothing. Retrying is safe
400 MALFORMED_JSON	Not valid JSON	Fix and resend. Do not retry unchanged
400 MISSING_FIELD	One of the five is missing or malformed	Fix and resend
400 INVALID_IDENTIFIER	An id has disallowed characters	Fix and resend
401 INVALID_KEY	Key missing or wrong	Stop and contact AIMS LAB
413 TOO_LARGE	Body over 1 MB	Check for a file or image in the JSON
422 SCHEMA_INVALID	Stored in quarantine , with the bad fields named	Not lost. Read the list, fix, resend
500 / 503	Our server has a problem	Retry with growing delay. Never drop the message

No rate limit applies to CMED. Seven clinics at fourteen rooms produce roughly one message every two seconds at peak. Send as consultations happen — but from a queue.

16 · Never let a message delay the doctor

Put both endpoints on a durable background queue. A slow network must never make Build Prescription feel slow. If our server is briefly unreachable the consultation still happened and the recording still exists — the clinical record just needs to catch up.

API 2 arrives...	What happens to the recording
within a second or two	Confirmed. Nobody notices anything
late — minutes, or an hour	Recording is never cut. It keeps going, we keep asking, and it is confirmed the moment API 2 arrives
never	Recording is still not cut. It is uploaded and marked <i>unconfirmed</i> , kept out of the dataset, and we are alerted. Deleted after 24 hours

A message arriving an hour late is completely fine. A message never sent is the only real failure. Notice what never happens: a missing message does not stop the doctor and does not cut a recording.

17 · The reply that looks like success but is not

	Proves	Does NOT prove
200 RECORDING_STARTED	The microphone is live and audio is being written	that the consultation will be kept
202 ACCEPTED	Your message is stored and will not be lost	that it matched a recording

There is a third step, on our server, a second or two later: **we compare the five fields the recorder reported against the five in your message, and the recording is only kept if they are identical.**

The failure that matters, stated plainly

If your `start_time` differs from API 1's by even a second, or loses its `+06:00`, or the clinic code came from your own settings instead of the laptop:

- the recorder says **200 RECORDING_STARTED**
- our server says **202 ACCEPTED**
- **CMED sees nothing wrong at all**
- and the recording is **deleted twenty-four hours later**

Two green replies mean both halves arrived. Whether they found each other is a separate question CMED cannot see the answer to. Section 22 is how to prove it.

Previous prescriptions

How history reaches us today, and the read API we will add

18 · Today: history comes to us inside API 2

There is no API for CMED to read previous prescriptions from us. The clinical surface is two POST endpoints and nothing else — data flows CMED to AIMS LAB only. Previous prescription information reaches us as the `previous_visit` object inside API 2.

The field must always be present. Send `null` when there is nothing to send. Leaving it out is reported as a problem, because we cannot tell "first visit" from "CMED forgot".

The situation	What to send
The patient's first ever visit	"previous_visit": null
A returning patient, and AIMS LAB may not hold the previous visit	the full previous visit, below
A returning patient, and AIMS LAB already holds it	"previous_visit": null

```
"previous_visit": {
  "date": "2026-06-02",          <-- REQUIRED if you send it
  "prescription": {
    "items": [
      { "drug": "Metformin 500 mg", "dose": "1 tablet",
        "frequency": "1+0+1", "duration": "30 days",
        "instructions": "after food" }
    ],
    "diagnoses": ["Type 2 diabetes mellitus"],
    "notes": "review in one month"
  }
}
```

Field	Rule
date	Required if you send <code>previous_visit</code> at all. YYYY-MM-DD, the date of that earlier consultation. Without it the whole previous visit is ignored
prescription	Same shape as API 3 Channel B, section 13. Only <code>items</code> really matters
diagnoses, notes	May sit inside prescription or beside it — both work

Sending it more often than necessary is completely safe. A repeated previous visit updates the one row we already have — nothing is duplicated. **The dangerous direction is the other one:** deciding not to send it when we do not actually hold it leaves a permanent hole in that patient's history that nobody notices.

Recommendation for the first release

Send `previous_visit` every time, for every returning patient. It costs about 2 KB per consultation. Because there is no read API yet, CMED has no way to know what we hold — so the only safe rule is to send it always. Switch to conditional sending once section 19 exists.

19 · The read API we will add, when you need it

You have said you will need to read previous prescriptions back in future. **This is not built yet.** Here is the design, so your team can plan around it and tell us if the shape is wrong before we write it.

19.1 The light one, which solves the conditional rule

PLANNED - not yet available

```
GET https://<aims-host>/api/v2/clinical/patient/{patient_id}/last-visit
X-CMED-Key: <key>
```

```
200 {
  "patient_id": "P0012345",
  "held": true,
  "last_visit_date": "2026-06-02",
  "prescription_versions": 2
}
```

```
200 { "patient_id": "P00999999", "held": false }    <-- we hold nothing
```

This is the one that matters most. One cheap call before sending API 2 tells CMED whether we already hold the previous visit, which turns the three-way rule in section 18 from a guess into a decision — and removes the need to send 2 KB of history on every single consultation.

19.2 The full one, for reading history back

PLANNED - not yet available

```
GET https://<aims-host>/api/v2/clinical/patient/{patient_id}/history?limit=5
X-CMED-Key: <key>
```

```
200 {
  "patient_id": "P0012345",
  "count": 2,
  "visits": [
    {
      "date": "2026-10-09",
      "doctor_id": "DR0042",
      "hospital_id": "AALO_DH0LPUR",
      "issued_at": "2026-10-09T10:26:55+06:00",
      "version": 1,
      "diagnoses": ["Acute gastritis"],
      "investigations": ["CBC", "HbA1c"],
      "advice": "avoid spicy food",
      "follow_up": "2026-11-09",
      "items": [
        { "line_no": 1, "drug": "Napa 500 mg", "dose": "1 tablet",
          "frequency": "1+1+1", "duration": "5 days",
          "instructions": "after food" }
      ]
    },
    { "date": "2026-06-02", "...": "..." }
  ]
}
```

Parameter	Meaning	Default
limit	how many visits back, newest first	5, max 50
from, to	restrict to a date range	none
hospital_id	only visits at one clinic	all

What it will and will not return

Will: the prescriptions, diagnoses, investigations, advice, follow-up dates and medicine lines CMED itself sent us, newest first, with their version numbers.

Will not: any audio, ever, and no transcript. Those are not served to CMED under any circumstances, and no parameter will change that.

Will not: another clinic's patients — the key is scoped to CMED's own hospitals.

	Before we build it, we need from CMED	Why
1	Confirmation that the shape above is what you want	Cheap to change now, expensive later
2	Whether you need 19.1, 19.2, or both	19.1 is a day's work; 19.2 is a few days. If the conditional rule is the only need, 19.1 alone is enough
3	Roughly how often you would call it	A per-consultation call and a nightly batch are sized differently
4	Whether a doctor's own notes should be included	Those are clinical text we hold; it is a policy question, not a technical one

Until this exists, section 18's recommendation stands: send previous_visit every time. It is the only rule that cannot lose history.

The data model

Where every field lands, and what the two databases hold

20 · Where every field ends up

So CMED can see that nothing is discarded. All of this lives in a database called `aims_clinical`, separate from the one that tracks recordings — which deliberately holds no patient names at all.

You send	We store it in	Holding
the whole message, untouched	<code>intake_records</code>	exactly what you sent , once per message, before anything is parsed
demographics	<code>patients</code> , <code>encounter_demographics</code>	name, sex, date of birth, phone, address, and the age at this visit
paramedic	<code>paramedic_observations</code>	each reading in its own column, blood pressure also split into systolic and diastolic, anything unrecognised in other
the visit itself	<code>encounters</code>	one row per consultation, linked to its recording
previous_visit	<code>encounters</code> , marked as a previous visit	a second row, dated to that earlier day
items	<code>prescription_items</code>	one row per medicine , in the order you sent them
the prescription	<code>prescriptions</code>	with a version number, so corrections never overwrite
diagnoses	<code>diagnoses</code>	one row each
investigations	<code>investigations</code>	one row each
female_details / male_details	their own tables	stored as sent

21 · The two databases, and why they are separate

	<code>aims_clinical</code>	<code>aims_recordings</code>
Holds	Everything CMED sends — patients, demographics, paramedic readings, prescriptions, previous visits	The audio tracking record — sessions, segments, hashes, archive state
Patient names?	yes	no — by design, none at all
Size	about 76 KB per consultation; roughly 10 GB for the study	small — metadata only
Who writes it	the two endpoints in this document	the recorders

The split is deliberate. A recording can be tracked, verified, archived and audited end to end without ever touching a patient name, which means the operational half of the system carries no identifying data.

A copy of every consultation's clinical record is **also written as a JSON file beside the audio**, with the same name. So each consultation is one folder containing the audio and its record together, and the database is an index into them — if the database were lost entirely, the dataset would still be readable.

Testing and handover

What to run, what to send us, and the page to keep on the desk

22 · Testing, in the order that makes sense

	What	Needs	How
1	The browser half — all of it	a PC with the recorder installed. No internet	Available today. Connect, doctors, API 1, the gate, API 3 Channel A
2	Check a server message against the rules	nothing at all	<code>python tools/channel_b_test.py --local</code>
3	The server half against a live host	a staging hostname from us	<code>channel_b_test.py --server https://<host> --key <key></code>
4	Watch every message and reply side by side	the staging site	<code>open /protocol-test</code>
5	A real consultation, all four messages	a clinic PC	and then step 6 — do not skip it
6	AIMS LAB confirms it matched and archived	ask us	the only check that proves the integration

Step 1 is unblocked today and is the half with the awkward detail — the gate, the session id, the provisional reply, reading the clinic code. Start there while we publish a staging host for step 3. **Please run step 3 and send us the output before the first clinic opens**; it takes under a minute and has caught every integration problem we have seen.

Step 6 must be asked for explicitly. Steps 1 to 5 can all pass while the five fields silently fail to match. We can see the answer on the dashboard in seconds — it is the only green light that means anything.

22A · Four commands that test the server half, before you write code

These need no application and no page. Run them as soon as you have the staging hostname and the key.

Is it up?

```
curl -s https://<aims-host>/health | jq .
# expect: "status": "healthy" AND "system": "AIMScribe v3"
```

Is my key accepted? Send a patient

```
curl -s -o /dev/null -w "%{http_code}\n" \
-X POST https://<aims-host>/api/v2/clinical/patient-information \
-H "X-CMED-Key: $CMED_KEY" \
-H "Content-Type: application/json" \
-d '{
  "patient_id": "TEST0001",
  "doctor_id": "DR0042",
  "hospital_id": "AALO_DHOLPUR",
  "start_time": "2026-10-09T10:14:32+06:00",
  "date": "2026-10-09",
  "demographics": { "name": "Test Patient", "sex": "female",
                    "age_years": 34 },
  "previous_visit": null
}'
# expect: 202
```

Send that exact command twice. Both times you get **202** — the second is **ALREADY_RECEIVED**. That proves retries are safe, which is the most useful thing to know before writing your queue.

Does validation work? Break it deliberately

```
# previous_visit omitted on purpose, and sex is a code rather than a word
curl -s -X POST https://<aims-host>/api/v2/clinical/patient-information \
  -H "X-CMED-Key: $CMED_KEY" -H "Content-Type: application/json" \
  -d '{ "patient_id": "TEST0002", "doctor_id": "DR0042",
        "hospital_id": "AALO_DHOLPUR",
        "start_time": "2026-10-09T10:20:00+06:00", "date": "2026-10-09",
        "demographics": { "name": "Test", "sex": "F" } }' | jq .
# expect: 422 SCHEMA_INVALID, naming demographics.sex and previous_visit
```

The reply names the bad fields. Build against that — it is faster than reading this document, and the record is quarantined rather than lost, so nothing is wasted by experimenting.

Send the prescription

```
curl -s -o /dev/null -w "%{http_code}\n" \
  -X POST https://<aims-host>/api/v2/clinical/prescription \
  -H "X-CMED-Key: $CMED_KEY" -H "Content-Type: application/json" \
  -d '{
    "patient_id": "TEST0001",
    "doctor_id": "DR0042",
    "hospital_id": "AALO_DHOLPUR",
    "start_time": "2026-10-09T10:14:32+06:00",    <-- SAME as API 2
    "date": "2026-10-09",
    "issued_at": "2026-10-09T10:26:55+06:00",
    "items": [ { "drug": "Napa 500 mg", "dose": "1 tablet",
                  "frequency": "1+1+1", "duration": "5 days" } ],
    "diagnoses": ["Acute gastritis"],
    "investigations": []
  }'
# expect: 202
```

The one thing these four commands cannot prove

All four can pass while the integration is still broken, because **a 202 means "stored", not "matched to a recording"**. If start_time differs by a second between API 1 and API 2, or loses its offset, every reply above is still green and **the recording is deleted twenty-four hours later**.

So the real test needs a recorder running and our confirmation — steps 5 and 6 of section 22. Please do not treat four green curls as integration complete.

23 · A complete worked example

```
let socket, sessionId = null, currentVisit = null, CLINIC = null;

function connect() {
  socket = new WebSocket("ws://127.0.0.1:5050/ws");

  socket.onopen = () => {
    // Read the clinic code off THIS laptop. Never hardcode it.
    socket.send(JSON.stringify({ command: "doctors" }));
  };

  socket.onmessage = (e) => {
    const m = JSON.parse(e.data);

    if (m.event === "status") { showRecordingState(m); return; }

    if (m.command === "doctors") {
      CLINIC = m.data.hospital_id;          // may be null = not enrolled
      if (!CLINIC) return showProblem("This PC is not registered");
      showClinicName(CLINIC);              // so a wrong PC is obvious
      fillDoctorSelector(m.data.doctors);  // [] is normal - do not block
      return;
    }

    if (m.command === "start") {
      if (m.code === "RECORDING_STARTED" ||
          m.code === "RECORDING_PROVISIONAL") {
        sessionId = m.data.session_id;      // keep it - API 3a needs it
        showRecording(true);
      } else {
        showProblem(m.code);                // act on code, not message
      }
    }
  };

  socket.onclose = () => { showRecording(false); setTimeout(connect, 2000); };
  socket.onerror = () => showProblem("AIMScribe is not running on this PC");
}

// The doctor opens a patient.
// Build the five fields ONCE, here, and reuse this object everywhere.
function openPatient(patient, doctor) {
  currentVisit = {
    patient_id: patient.id,
    doctor_id:  doctor.id,
    hospital_id: CLINIC,                // from the recorder, not settings
    start_time: new Date().toISOString(),
    date:       localDateString()       // the clinic's local day
  };

  socket.send(JSON.stringify({
    command: "start",
    request_id: "open-" + currentVisit.patient_id,
    trigger: currentVisit               // the SAME object
  }));

  queueToOurBackend("API2", currentVisit, demographics, paramedic,
                    previousVisit);    // never inline - always queued
}

// The doctor presses Build Prescription. This does NOT stop recording.
function prescriptionBuilt(prescription) {
  socket.send(JSON.stringify({
    command: "prescription_built",
    request_id: "built-" + currentVisit.patient_id,
    patient_id: currentVisit.patient_id,
    session_id: sessionId
  }));

  // currentVisit is reused, so start_time is the ORIGINAL. This is the
  // single most important line in the integration.
  queueToOurBackend("API3B", currentVisit, prescription);
}
```

```
}  
connect();
```

24 · What each side must hand over

AIMS LAB gives CMED	When
The staging and production hostnames	with the key
The X-CMED-Key — shown once	on request, out of band
The seven clinic codes	with the key
channel_b_test.py	with the key
Confirmation a test recording matched and archived	on request, step 6
CMED gives AIMS LAB — both blocking	Why it blocks
1. Every exact web address your page is served from, production, staging and preview	Without it the page cannot open the local connection at all. Nothing works. www and non-www differ, as do http and https and any non-default port
2. Your own clinic codes, so we can map them to ours	Without the mapping a visit cannot be tied to the clinic it happened in. The record is stored, you receive a normal 202, and the recording is erased after 24 hours. The failure is completely silent from CMED's side
A named technical contact	So a question either way does not wait a week
A staging environment we can send to	So we can test our side against yours

25 · One page for the CMED developer

Question	Answer
How many messages do I build?	Four. Two from the browser, two from your server
What ties them together?	The five fields, identical in all four
Biggest trap	start_time in API 3 is the original, not when the prescription was built
Where do I connect, browser side?	ws://127.0.0.1:5050/ws — the same PC, every time
Do I need a key there?	No. We register your page's web address instead
Where do I POST, server side?	/api/v2/clinical/patient-information and /api/v2/clinical/prescription
How do I authenticate there?	X-CMED-Key. Server to server only, never in a browser
What do I keep from API 1?	session_id — API 3 Channel A needs it
Where does hospital_id come from?	The doctors command. Never hardcoded
Is RECORDING_PROVISIONAL bad?	No — it means recording. But keep listening; it can still be refused
Does prescription_built stop it?	No. It arms the gate; the next patient closes the consultation
Can I send stop or pause?	No. Those belong to the doctor
Required in API 2?	demographics, and previous_visit present (null is fine)
Required in API 3 Channel B?	issued_at, items, diagnoses, investigations — lists may be empty
Out-of-range reading?	Stored empty. It never rejects the message
Extra fields?	Kept. Never an error
Is retrying safe?	Yes. One row, and ALREADY_RECEIVED
422?	Nothing lost. Quarantined, with the bad fields named

Can I read previous prescriptions back?	Not yet — no read API exists. Send <code>previous_visit</code> every time. Section 19 is the planned design
Does a 202 mean it worked?	No — stored, not matched. Section 17
What must I send AIMS LAB?	Your page's exact addresses, and your clinic codes. Both blocking
Can I start today?	Yes — the whole browser half. Section 22

Prepared by AIMS LAB, United International University. **Document 1 of 2**; the companion is "Server Deployment", for the team helping us deploy. Every rule here is taken from the code that handles these messages. If anything is unclear, ask before building — it is far cheaper than discovering it in a clinic.